

MEMO

Never forget a word.

Your AI-powered meeting intelligence agent.

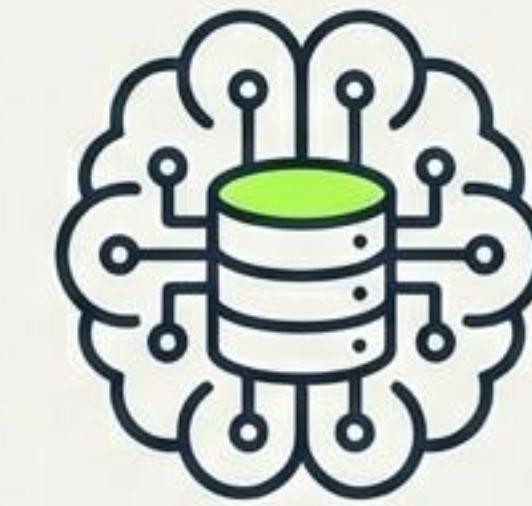
Meeting Amnesia

- ✖ Critical context is lost.
- ✖ Searching words in a video is impossible.
- ✖ Current solutions have no extensibility.

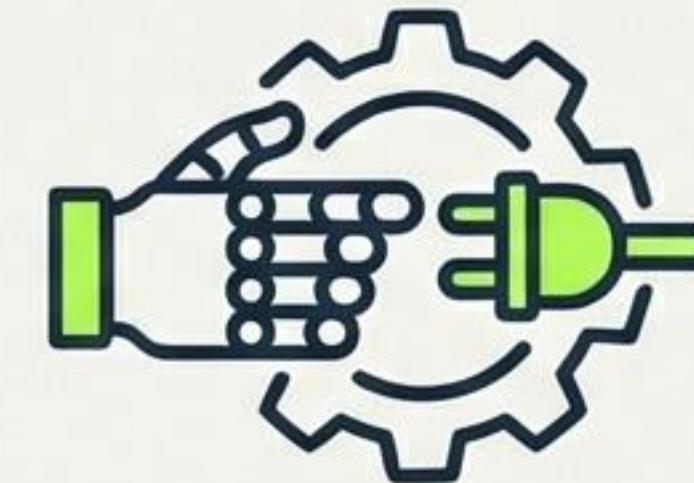
A Multimodal Agent



LISTENS



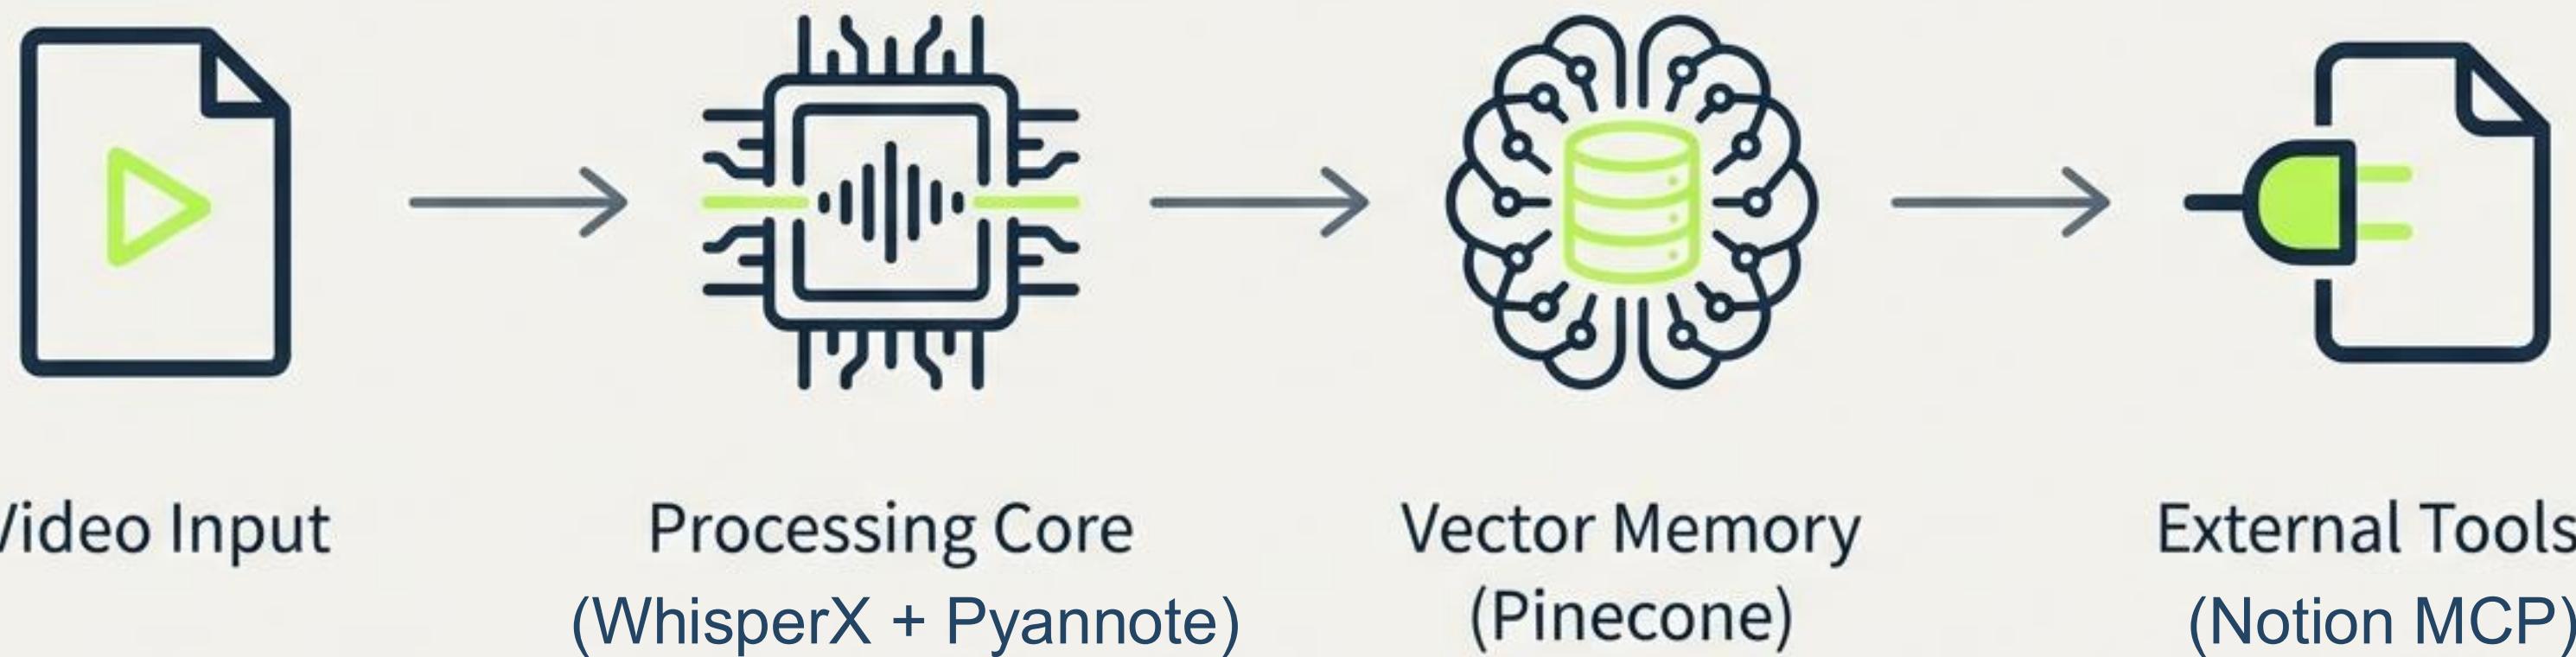
REMEMBERS



ACTS

High-Level Architecture

Video In → Intelligence Out.



Video Input

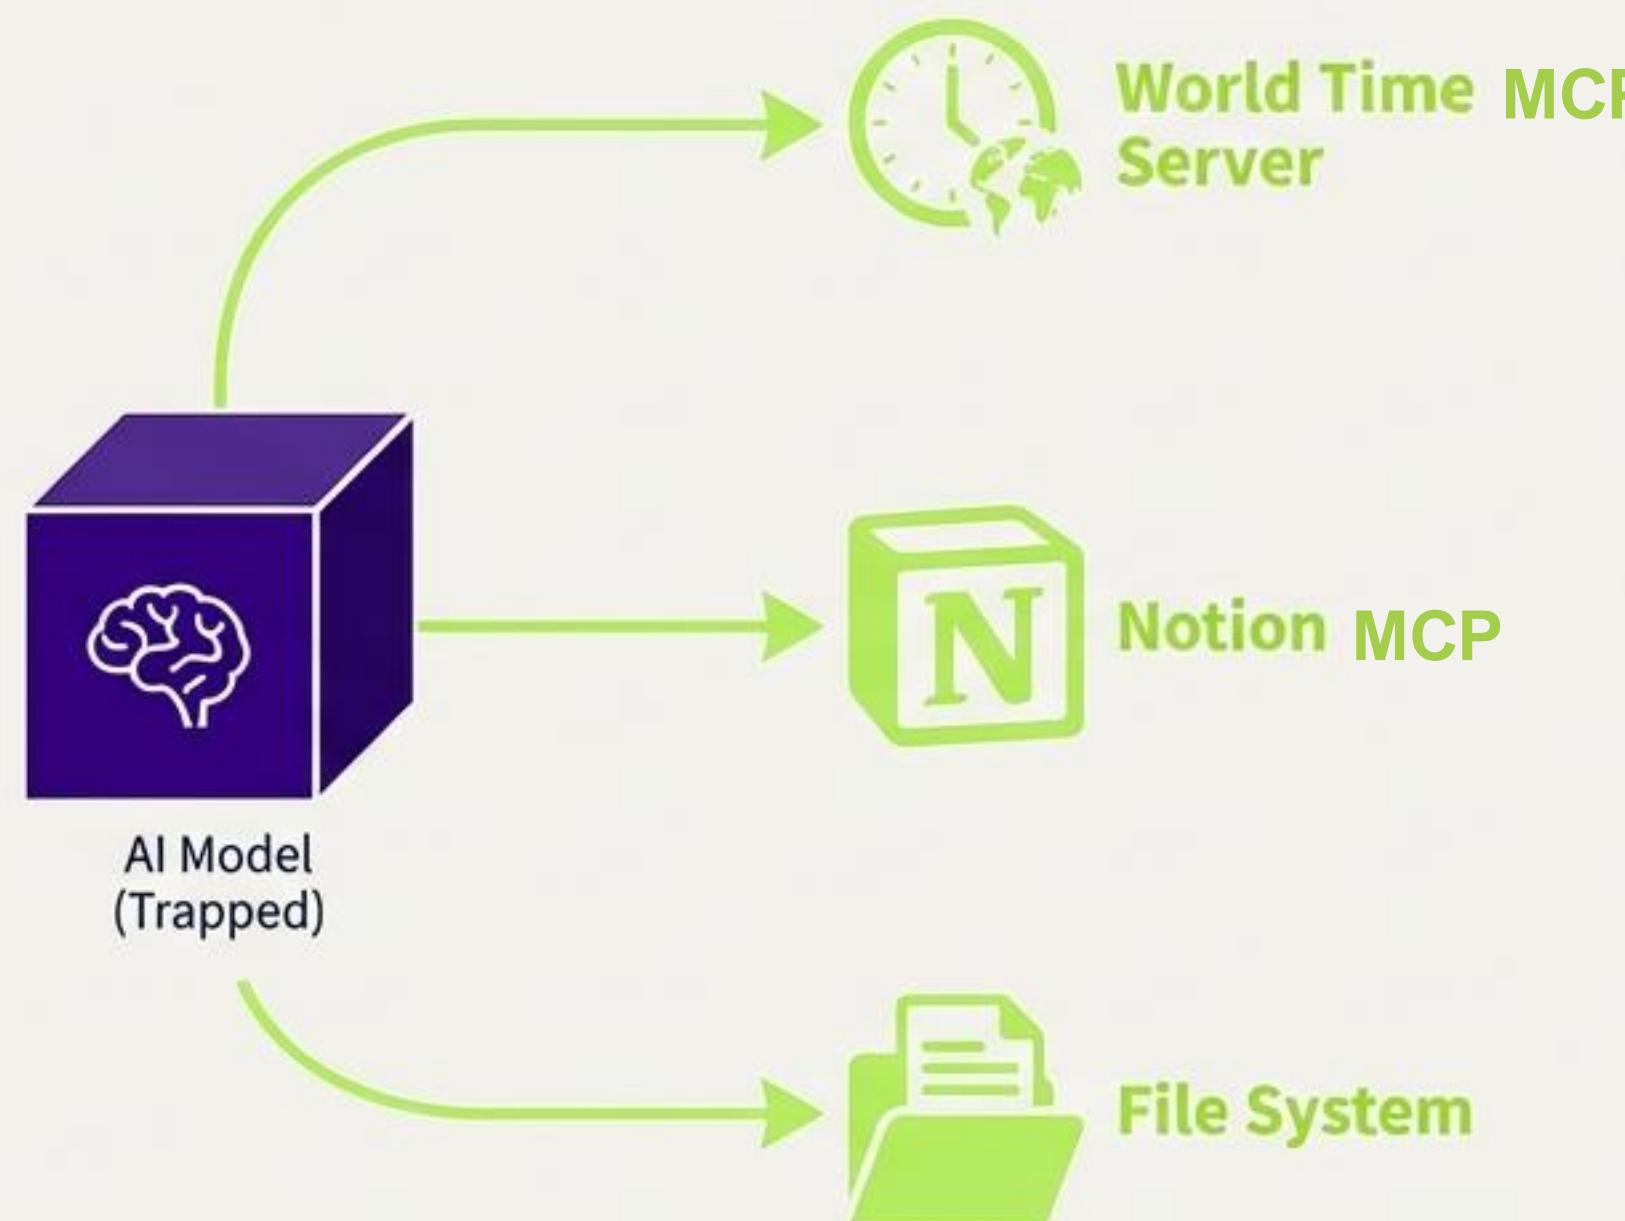
Processing Core
(WhisperX + Pyannote)

Vector Memory
(Pinecone)

External Tools
(Notion MCP)

MCP: The Universal Adapter for AI

Standard AI is trapped in a digital box. The Model Context Protocol (MCP) gives it “hands” to interact with real-world data and tools.

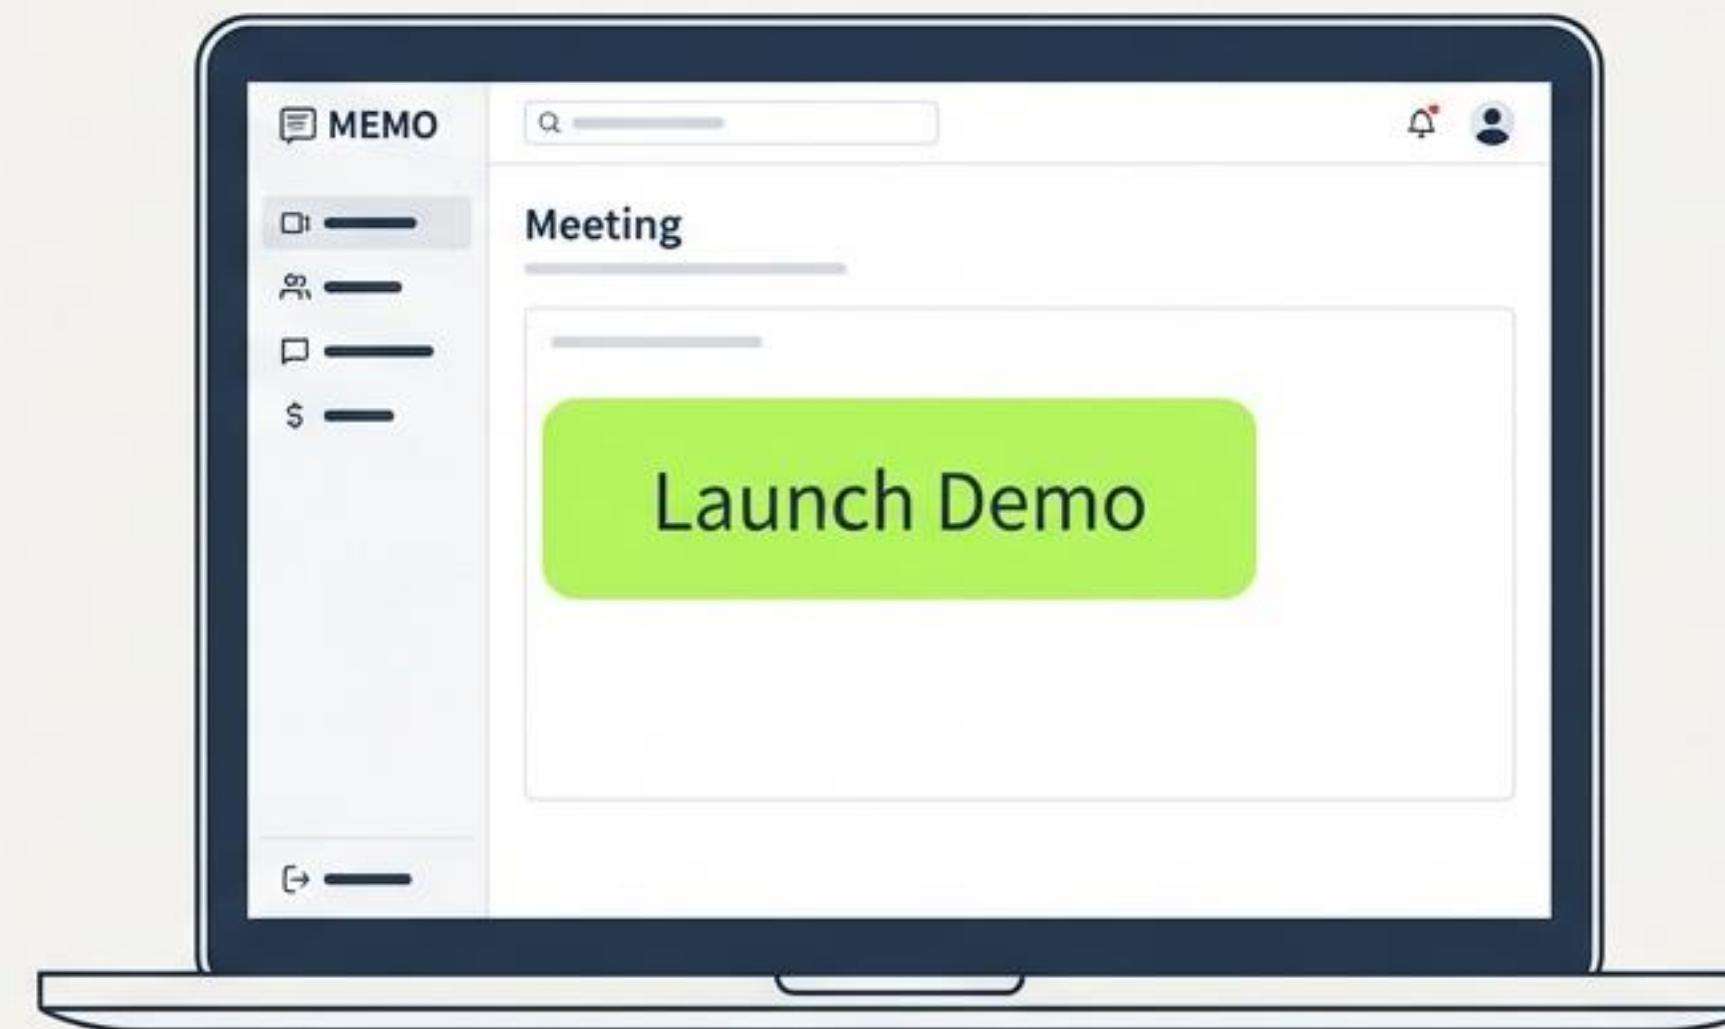


Enables **Time-Aware Queries**.
The agent can understand relative time by connecting to a custom World Time Server.

Query: “What did we decide about the budget last Tuesday?”

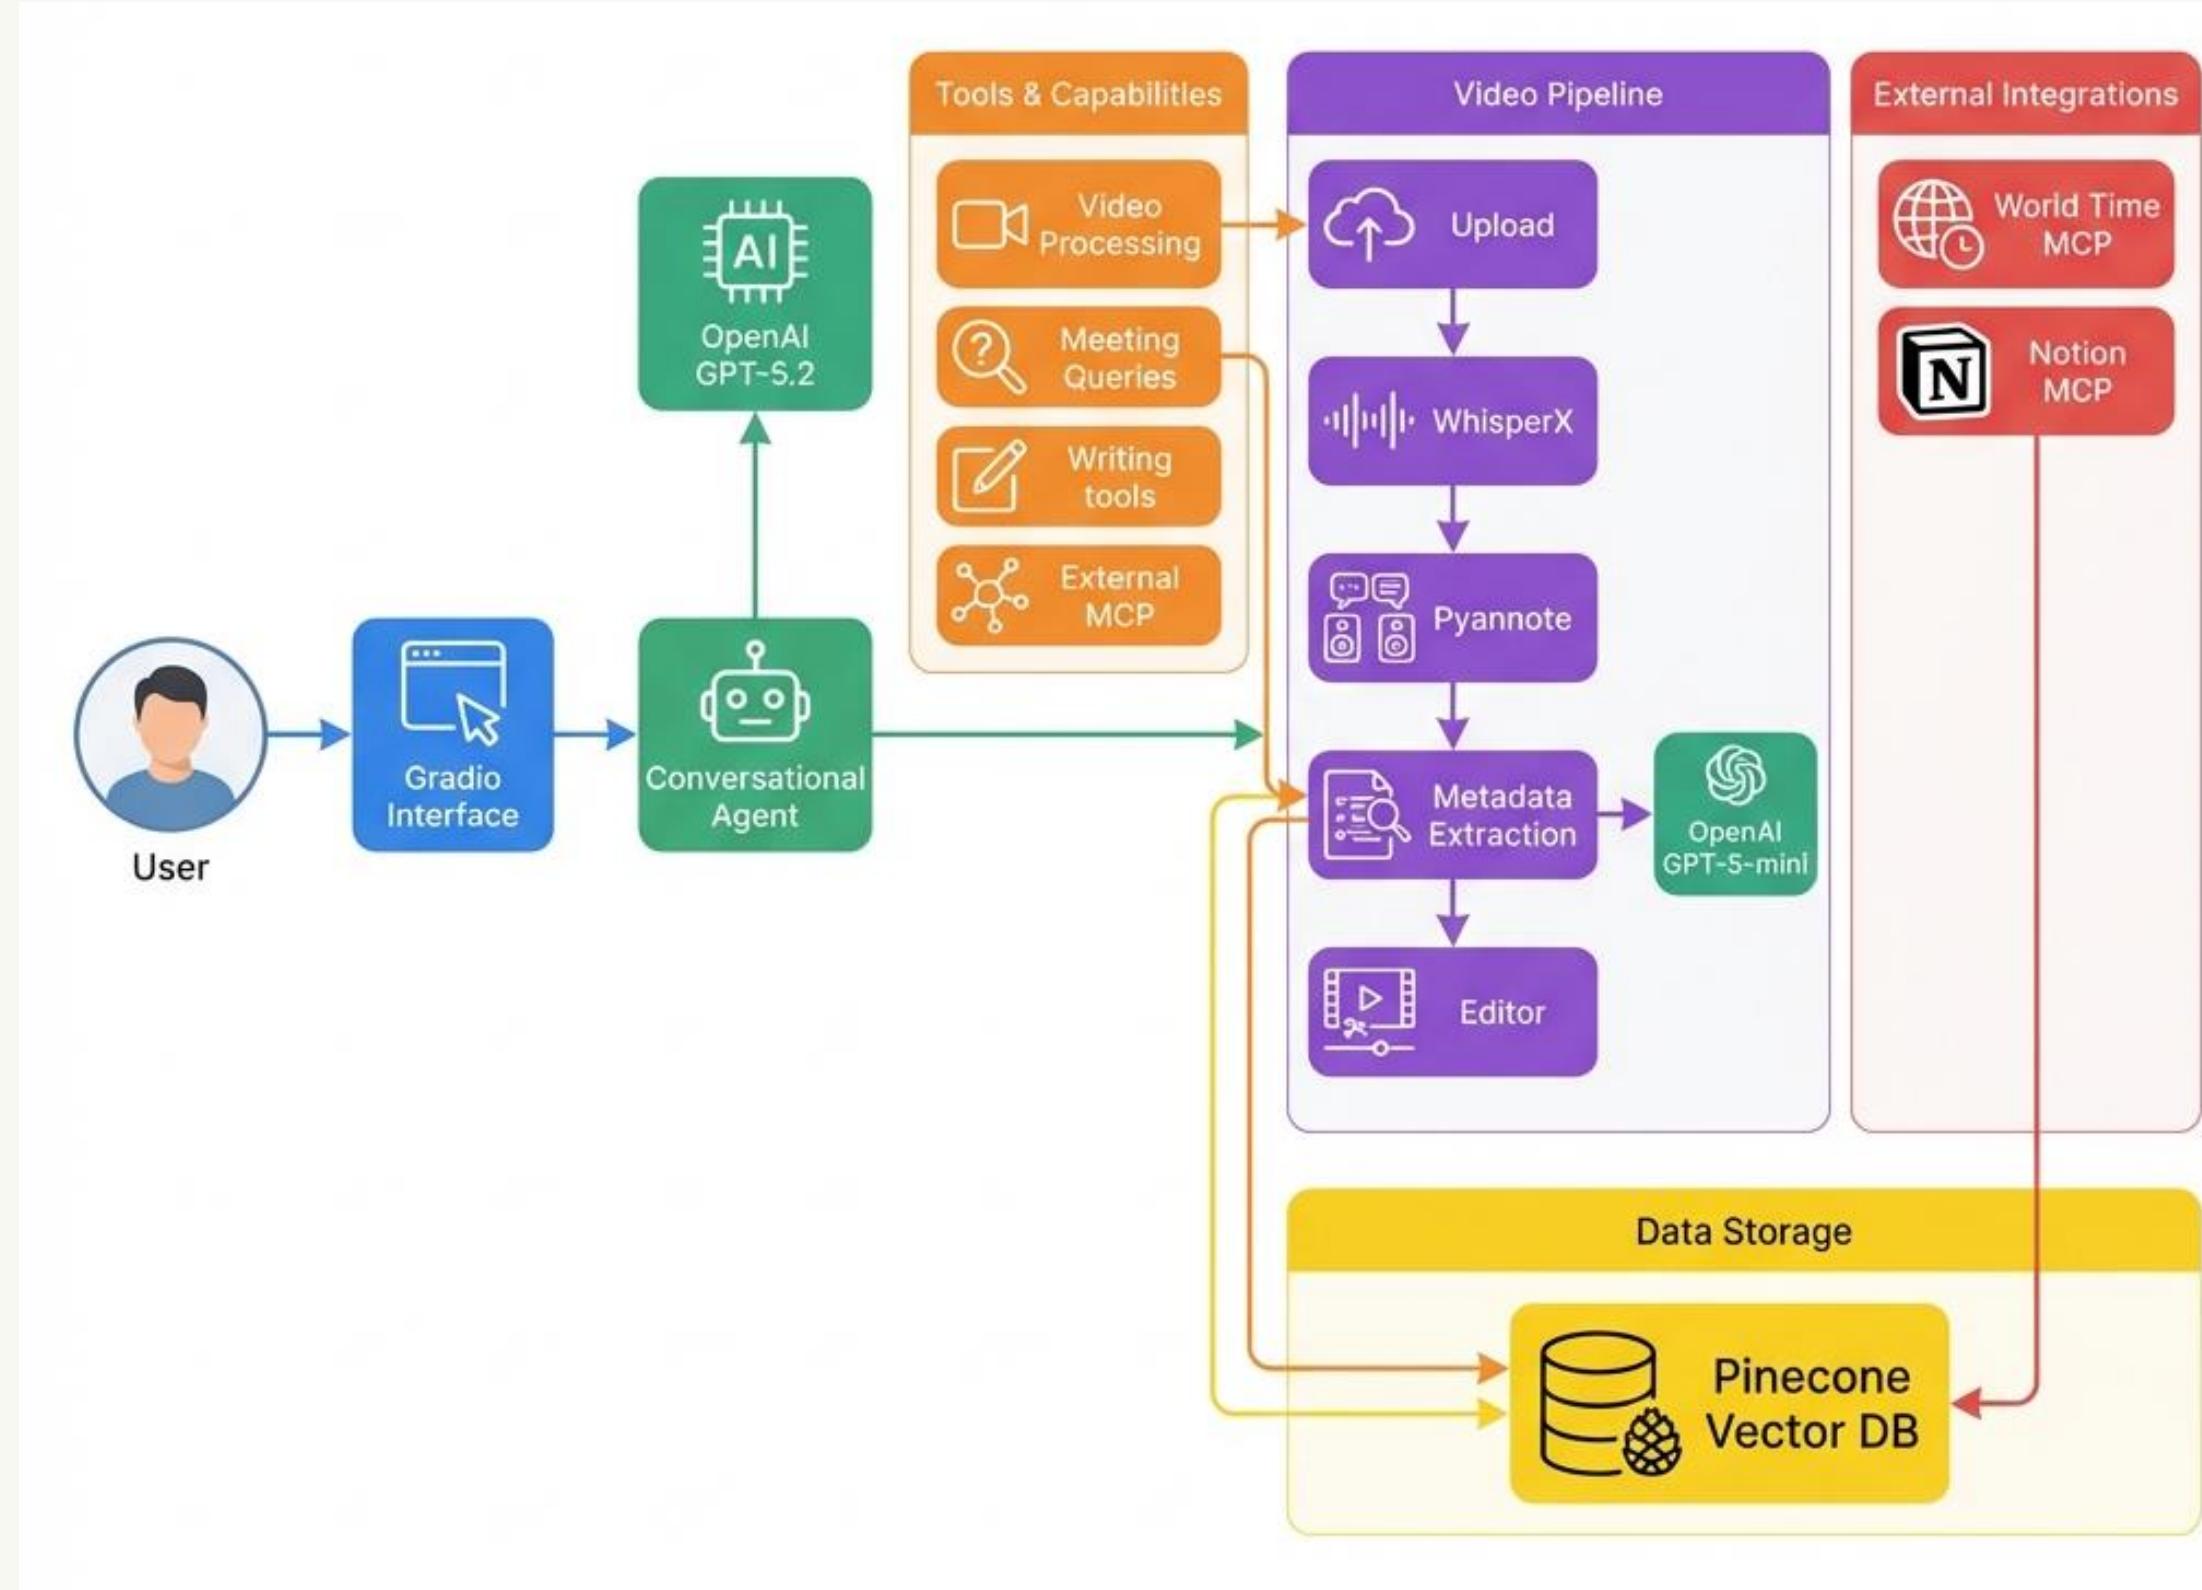
Seeing Is Believing

Explore the live application on Hugging Face Spaces.



<https://huggingface.co/spaces/GFiaMon/meeting-agent-docker>

Under the Hood: The Core Tech Stack





MEMO

Never forget a word.

An end-to-end, open-source
meeting intelligence solution.

Connect with me!



linkedin.com/in/guillermo-fiallo-montero-734a87132/

Roadmap: What's Next



Live Integration

Real-time stream processing
from Zoom & Teams meetings
via the RTMS API.



Student Mode

A specialized function to
automatically “summarize this
lecture for my exam,”
identifying key concepts and
definitions.



Agentic Actions

Proactively draft and send
follow-up emails with assigned
action items directly to
attendees.